



WPI

CS 4341

Introduction to Artificial Intelligence

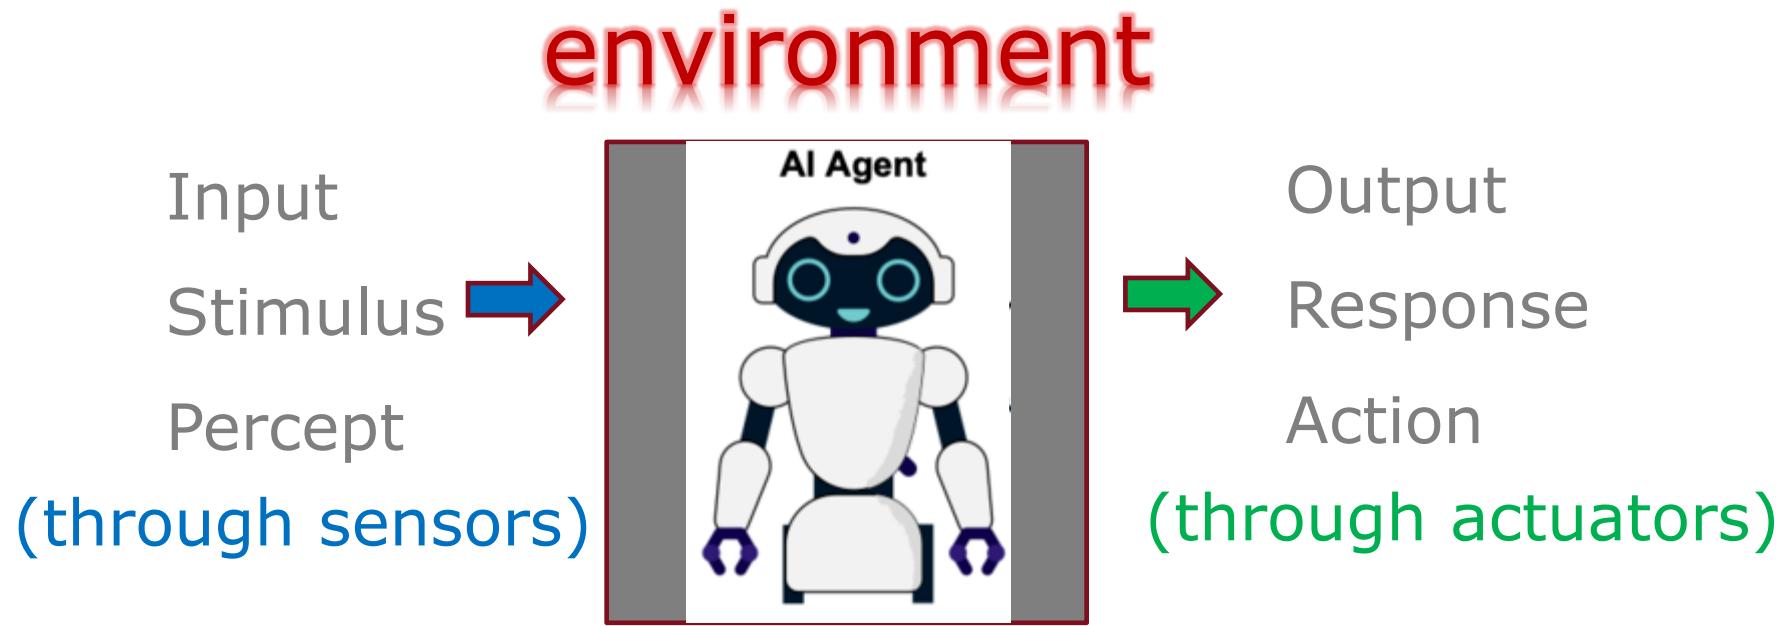
Lecture 2 and 3: Intelligent Agents

By

Ben C.K. Ngan

Worcester Polytechnic Institute

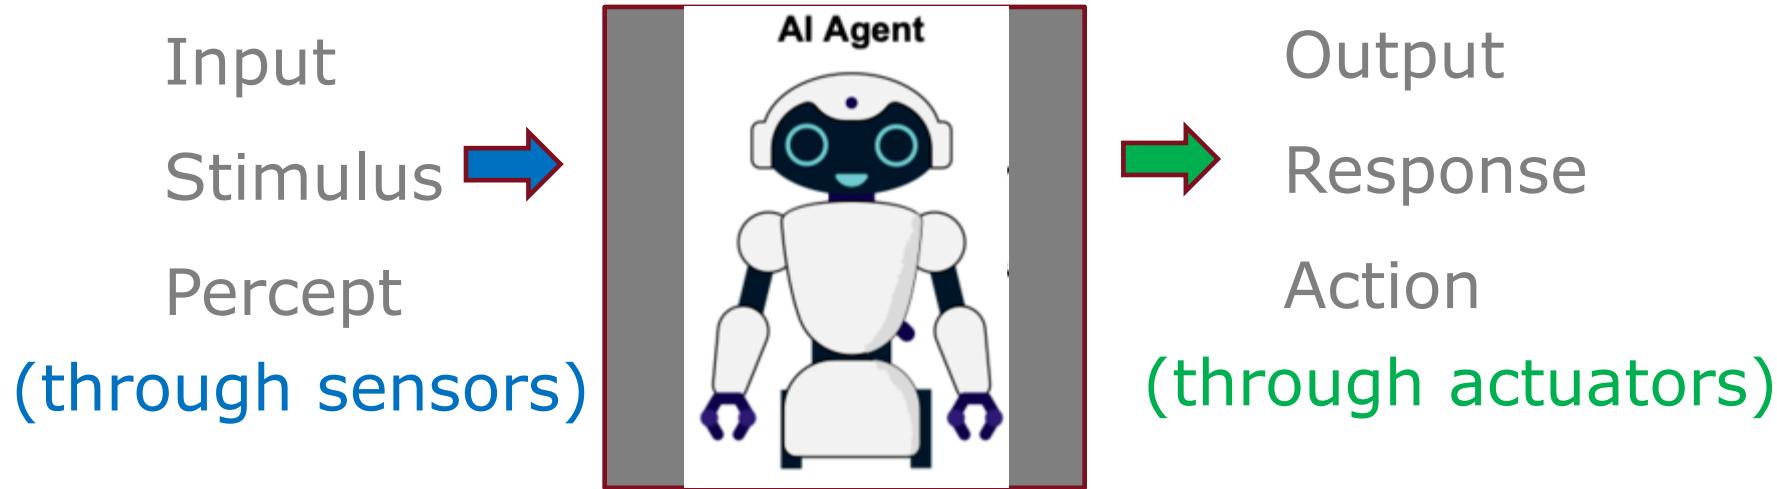
What is an Intelligent Agent?



- An **intelligent agent** (IA) is an entity (e.g., Robot, Chatbot, Self-driving Car, etc., to name a few) that includes a software program and a hardware architecture (sensors, mics, cameras, and actuators) to **perceive** (e.g., to see, smell, taste, hear, or touch) information from and **interact** with the environment, then **understand and think humanly/rationally/autonomously/adaptively** (i.e., performs computations), and finally decide to **take a sequence of decisions**, i.e., actions and responses, **humanly/rationally/autonomously/adaptively**, based upon the changes in its environment to maximize its **expected utility**, i.e., achieve a particular goal or optimize a particular objective function of the AI applications.
 - ❑ A loss function or a utility function is the **expected performance value** (e.g., accuracy, correctness, # of wins, etc.,) used to explicitly quantified each objective.
 - ❑ The **actual performance metric** (e.g., accuracy, correctness, # of wins, etc.,) is used to evaluate the effectiveness of IAs.

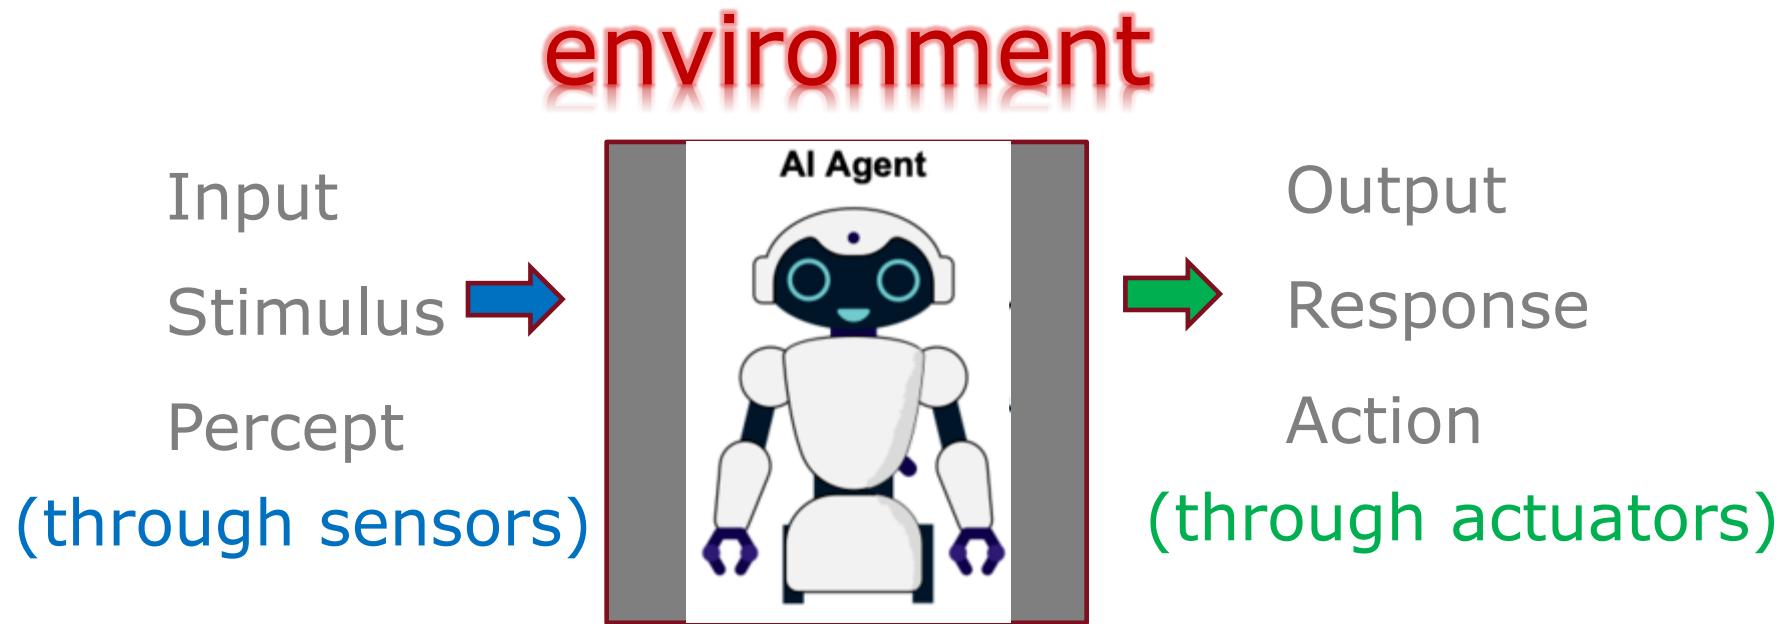
What is an Intelligent Agent?

environment



- The Intelligent Agent **NOT ONLY** interacts with the environment **BUT ALSO** interacts with other Intelligent Agents and then reacts to the changes in its environment.
- An Intelligent Agent, like a robot, exhibits several of the properties given below:
 - Autonomous: Proactively initiates the activities
 - Mobile: Move itself from the current point to another
 - Adaptable: Adapt to a new changing environment
 - Knowledgeable: Reason about the information and knowledge
 - Collaborative: Communicate with other intelligent agents and work in a cooperative manner
 - Persistent: Retain their knowledge and **state** over an extended period of time.

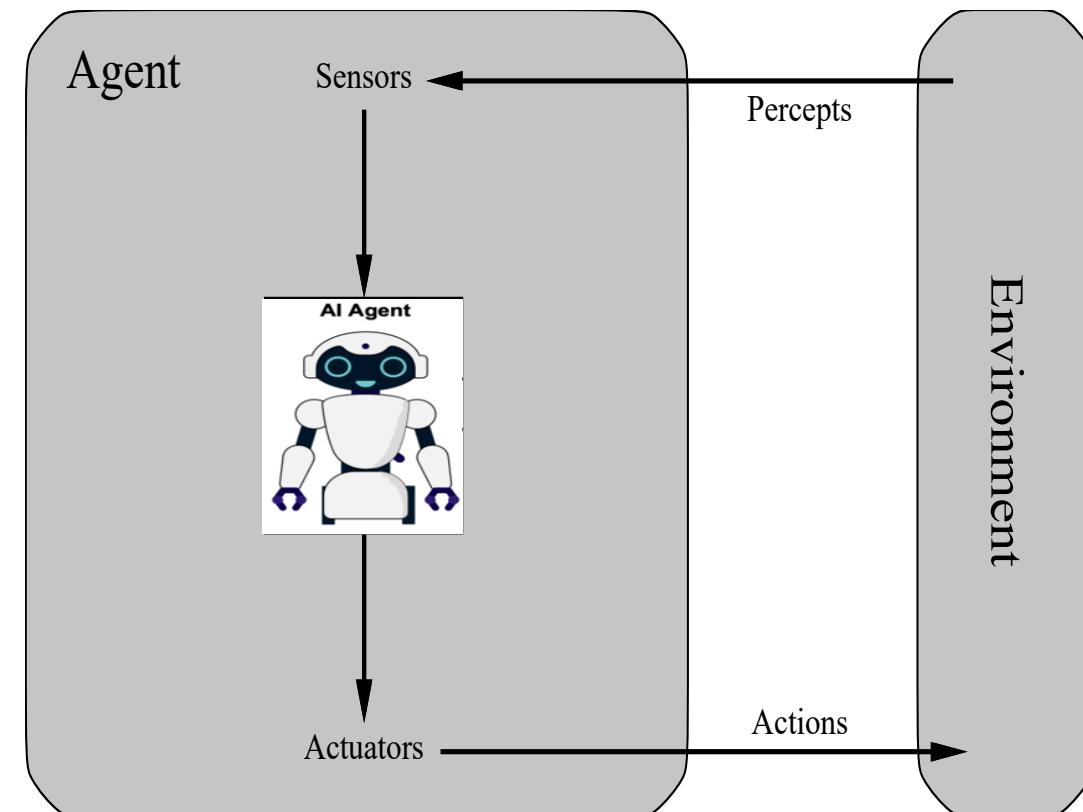
What is an Intelligent Agent?



- In the context of an AI application, a state corresponds to the internal/current configuration/situation of the all the variables and parameters of an intelligent agent that represents an environment to solve a problem at hand.
 - Some examples are the current position of the robot, the piece position on a chess-playing system, or the spatial location of the agent in a routing application.
 - Each next action of an agent leads to a transition to either the same state or a different state.
 - A transition may lead to a reward/a penalty, i.e., a utility, that directs the next actions of the agent.

What is an Intelligent Agent?

- Can **perceive** the state of the environment through its **sensors**
 - For humans, e.g., eyes, ears, noses, skins, etc.
 - For intelligent agents, e.g., motion cameras, mics, infrared range finders, etc.
- An intelligent agent's decision is described by the agent function that **maps** sensors (**any given percept sequence**) to actuators (**an action**), called the **control policy** for the agent.
- Can **act** upon its states through its **actuators**
 - For humans, e.g., hands, legs, mouths, vocal tract, etc.
 - For agents, e.g., motors, artificial limbs, speakers, etc.
- **Decisions** on which actuators to carry out based on the sensors are repeated in the **loop**
- The whole process is called **Perception-Effect Cycle**

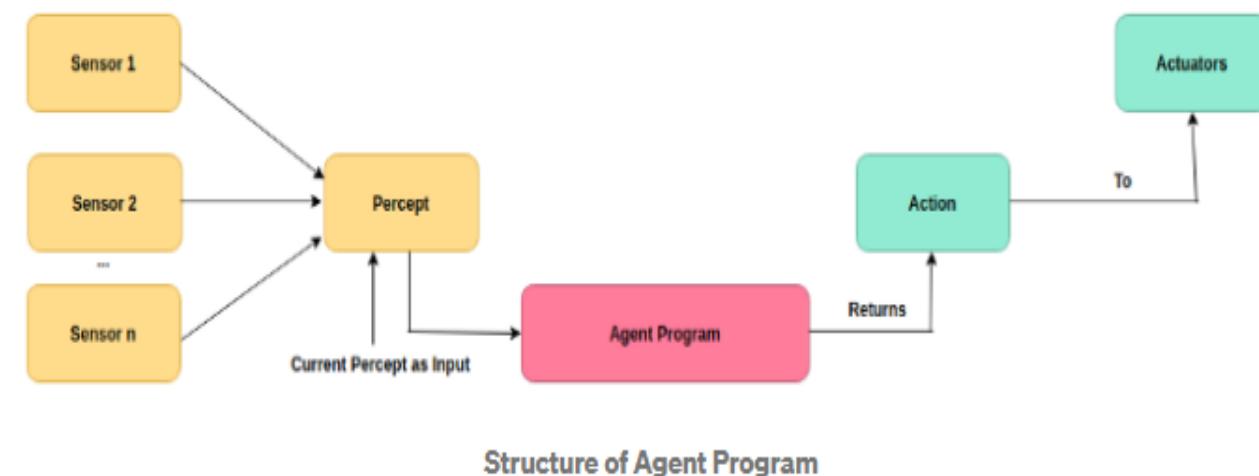


What is an Intelligent Agent?

environment



- Intelligent Agent = Software Program + Hardware Architecture
- This agent mapping / function is implemented using an agent program that runs on an architecture, i.e., hardware.
- The program runs on a computing device with physical sensors and actuators, i.e., an agent architecture.



What is an Intelligent Agent?



What is an Intelligent Agent?

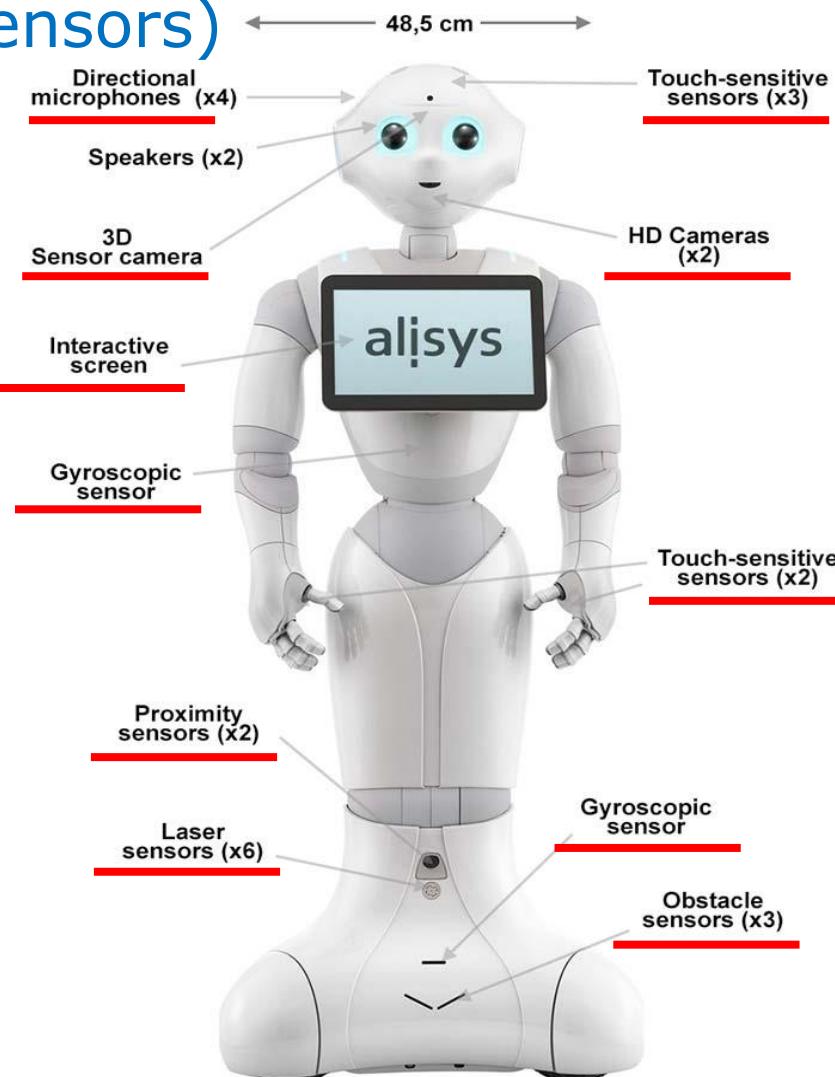


What is an Intelligent Agent?

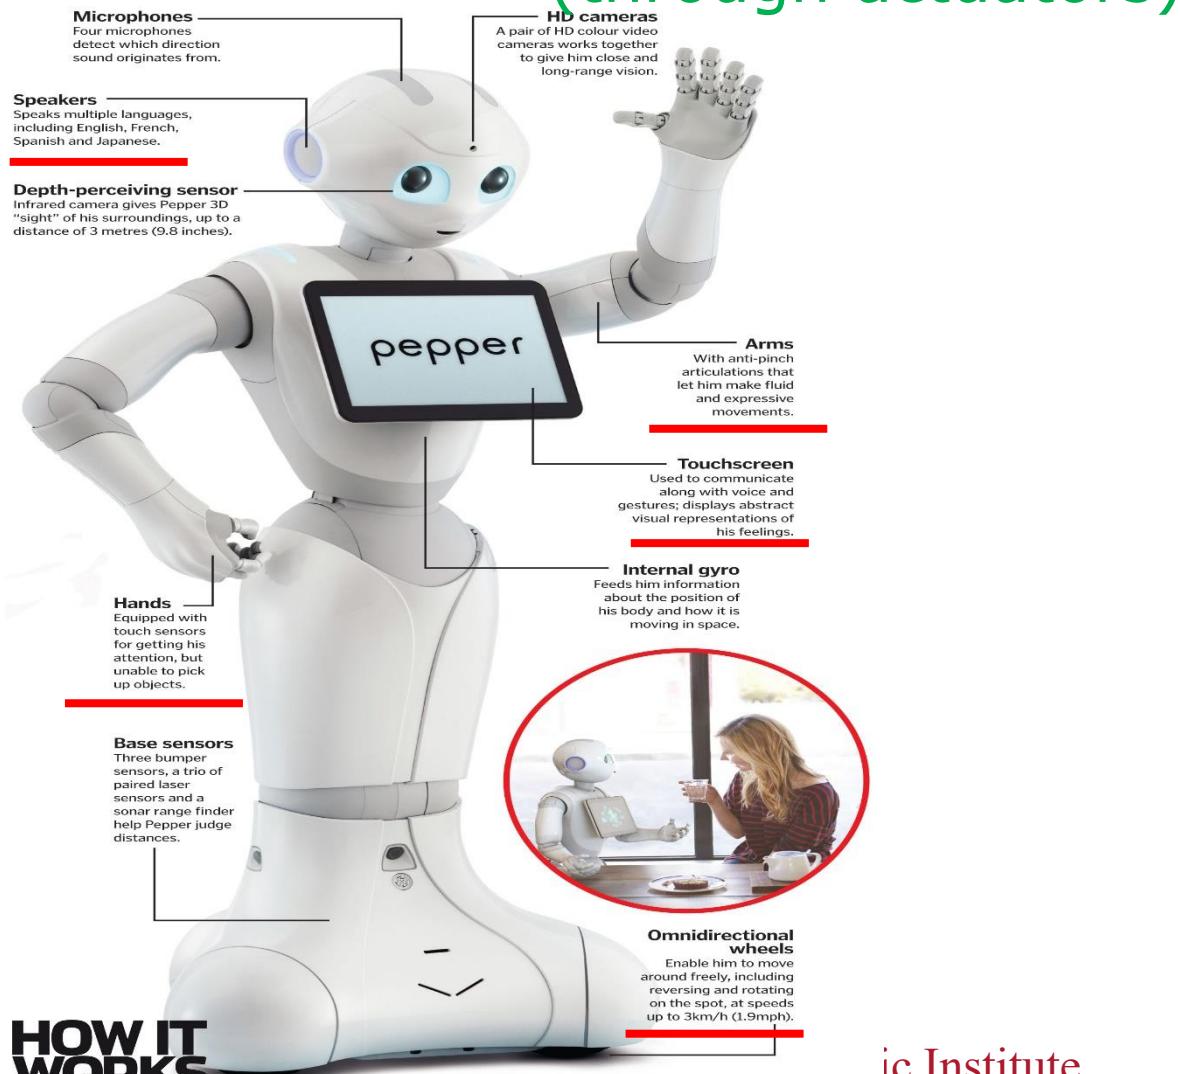


What is an Intelligent Agent?

(through sensors)



(through actuators)



What is an Intelligent Agent?



What is an Intelligent Agent?

Under the bonnet

How a self-driving car works

Signals from **GPS (global positioning system)** satellites are combined with readings from tachometers, altimeters and gyroscopes to provide more accurate positioning than is possible with GPS alone

Radar sensor

Lidar (light detection and ranging) sensors bounce pulses of light off the surroundings. These are analysed to identify lane markings and the edges of roads

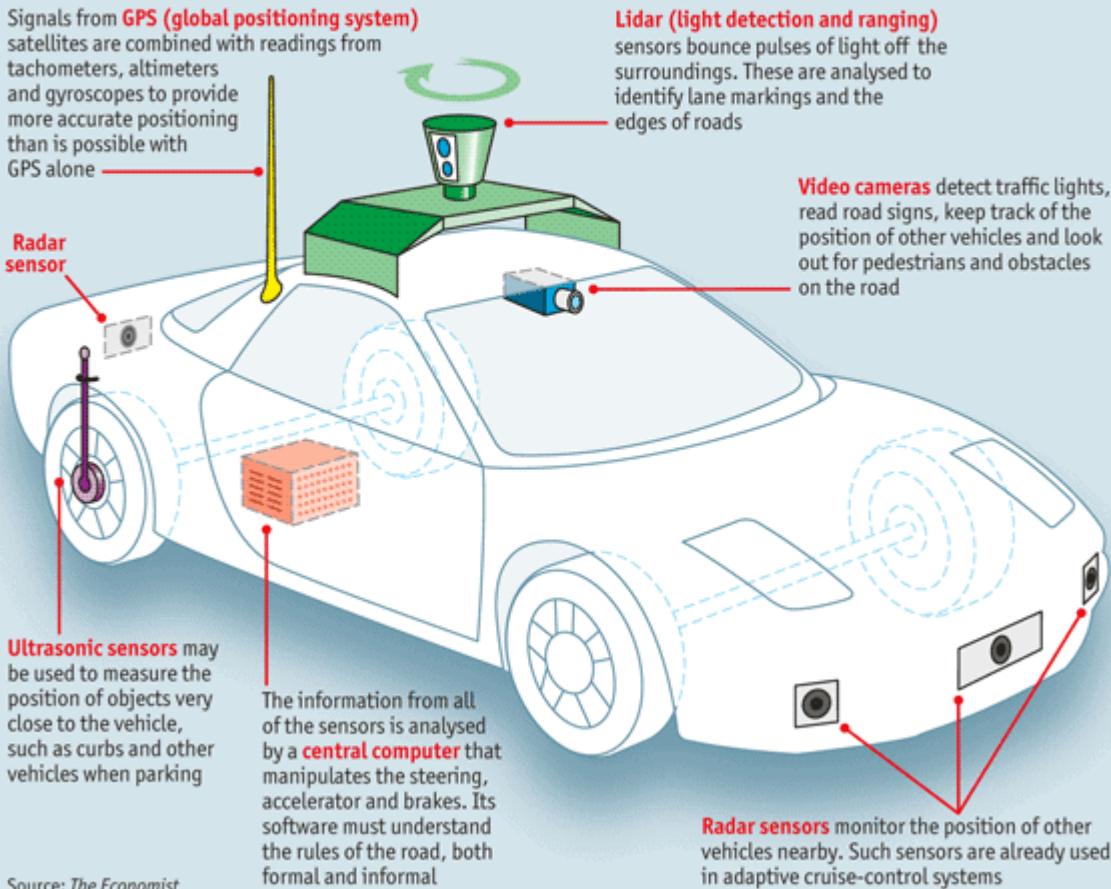
Video cameras detect traffic lights, read road signs, keep track of the position of other vehicles and look out for pedestrians and obstacles on the road

Ultrasonic sensors may be used to measure the position of objects very close to the vehicle, such as curbs and other vehicles when parking

The information from all of the sensors is analysed by a **central computer** that manipulates the steering, accelerator and brakes. Its software must understand the rules of the road, both formal and informal

Radar sensors monitor the position of other vehicles nearby. Such sensors are already used in adaptive cruise-control systems

Source: *The Economist*



Connectivity

Internet between vehicle-to-vehicle and vehicle-to-infrastructure system

Actuator

Take prompt action based on computer results

Mapping

Stores and updates geological and infrastructure information

Processors (ECUs/MCUs)

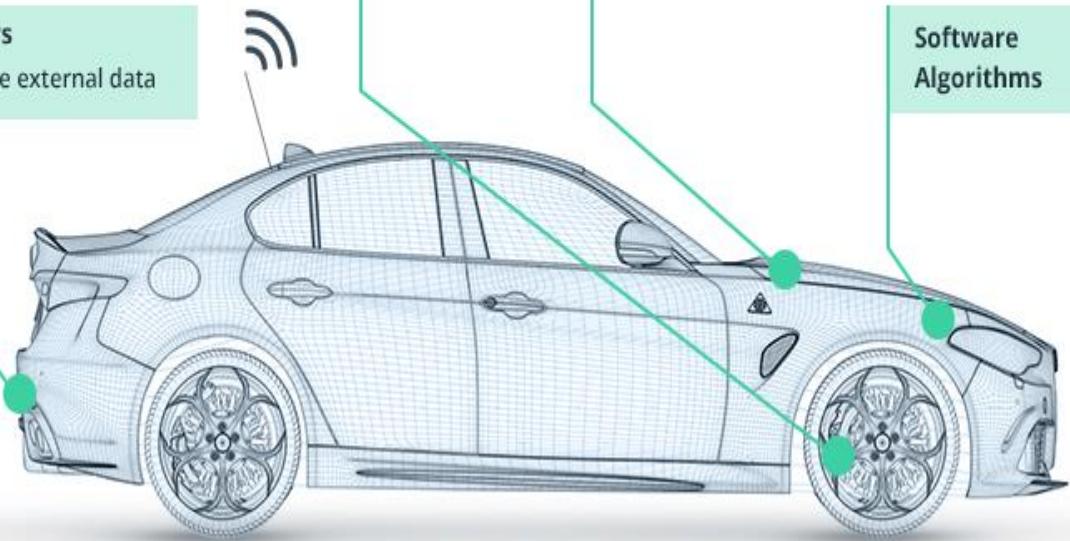
Process data needed to make decisions

Middleware

Software Algorithms

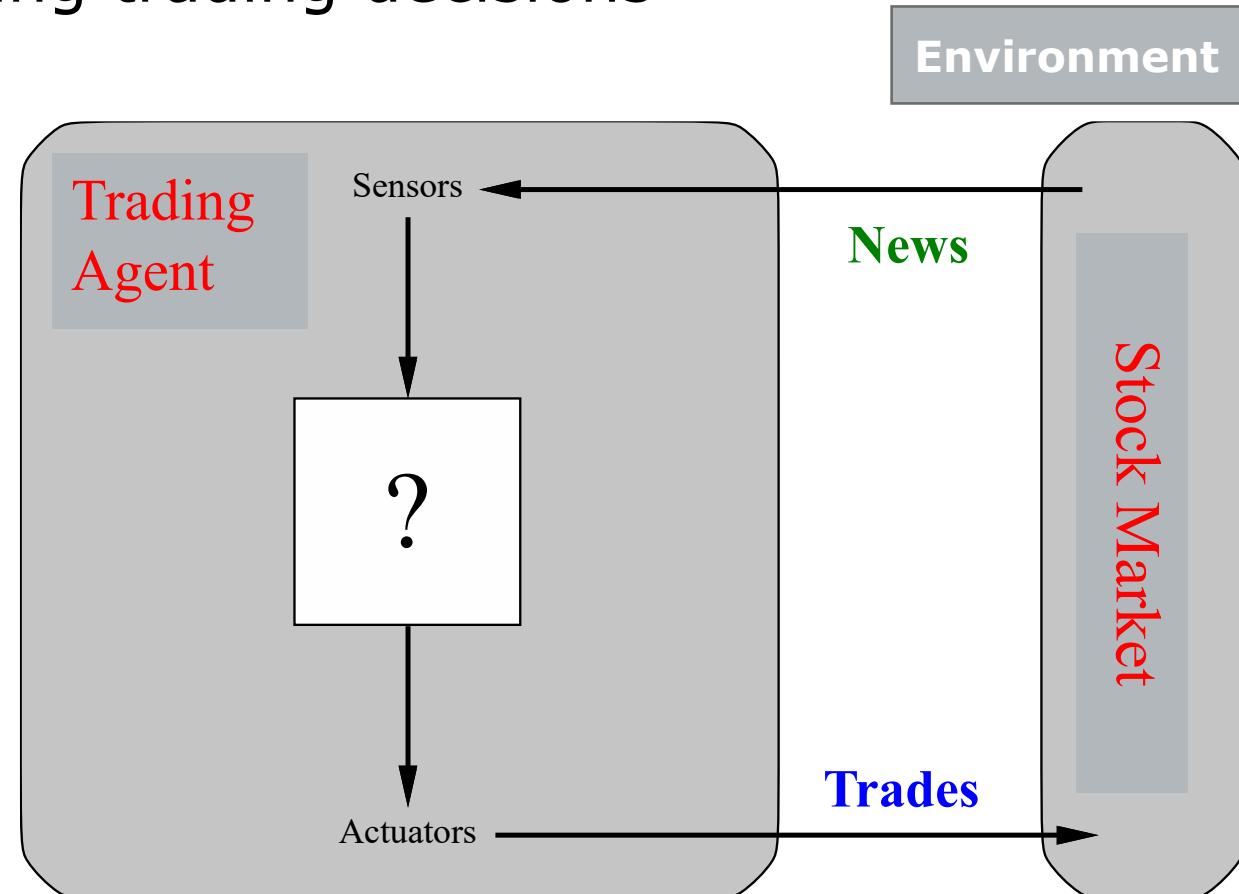
Sensors

Perceive external data



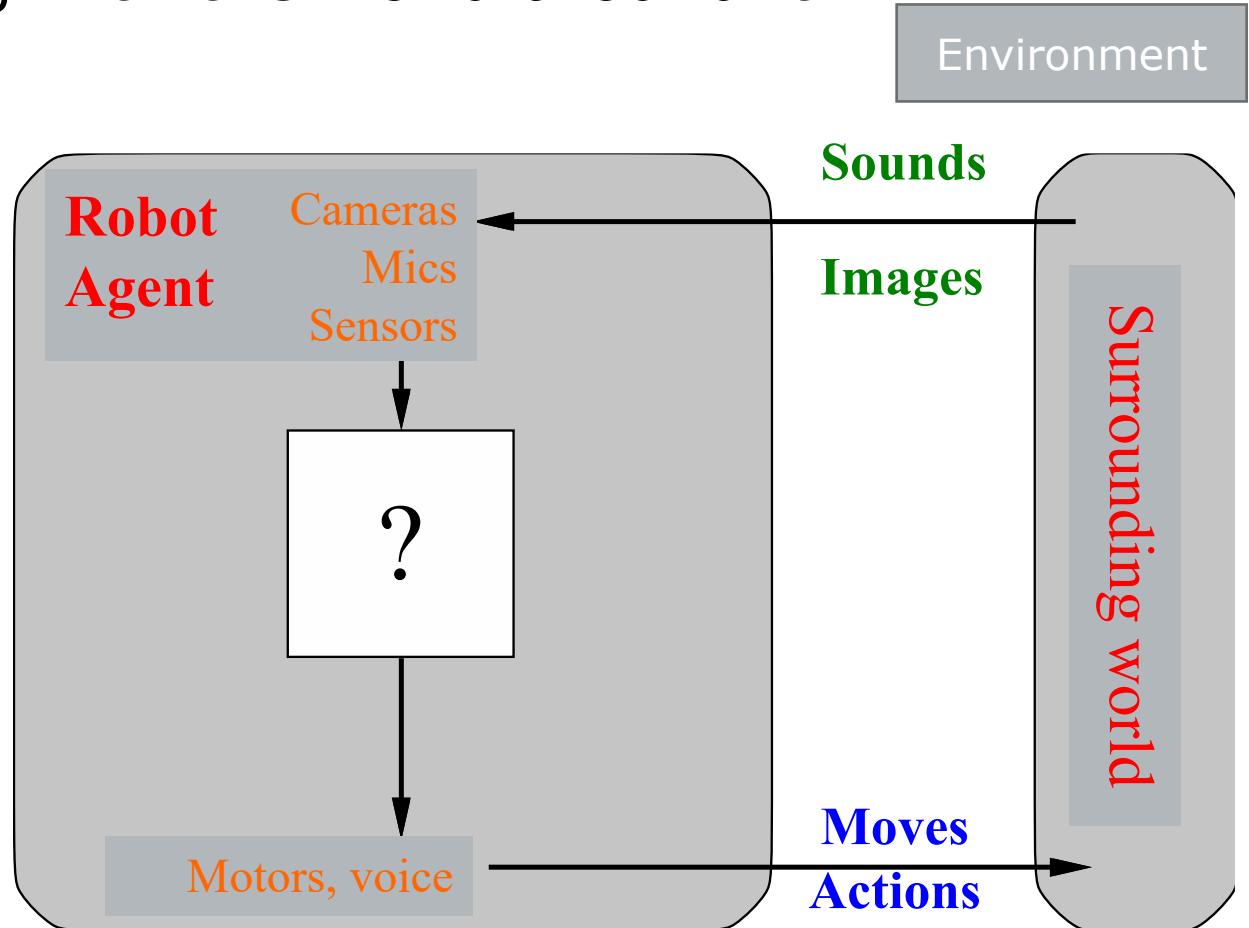
More Specific Example: AI in Finance

Trading Agent: Making trading decisions



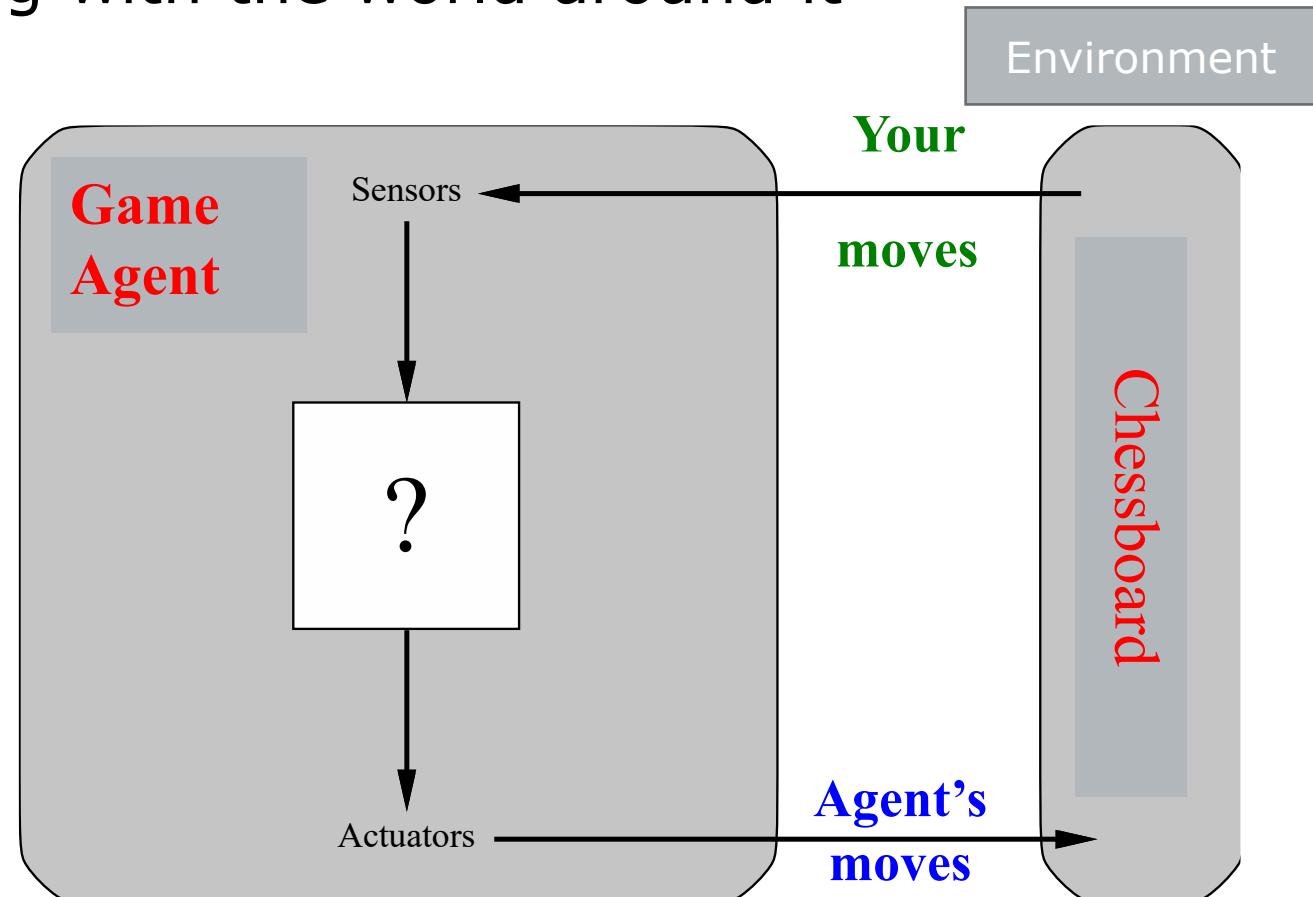
More Specific Example: AI in Robotics

Robot Agent: Interacting with the world around it

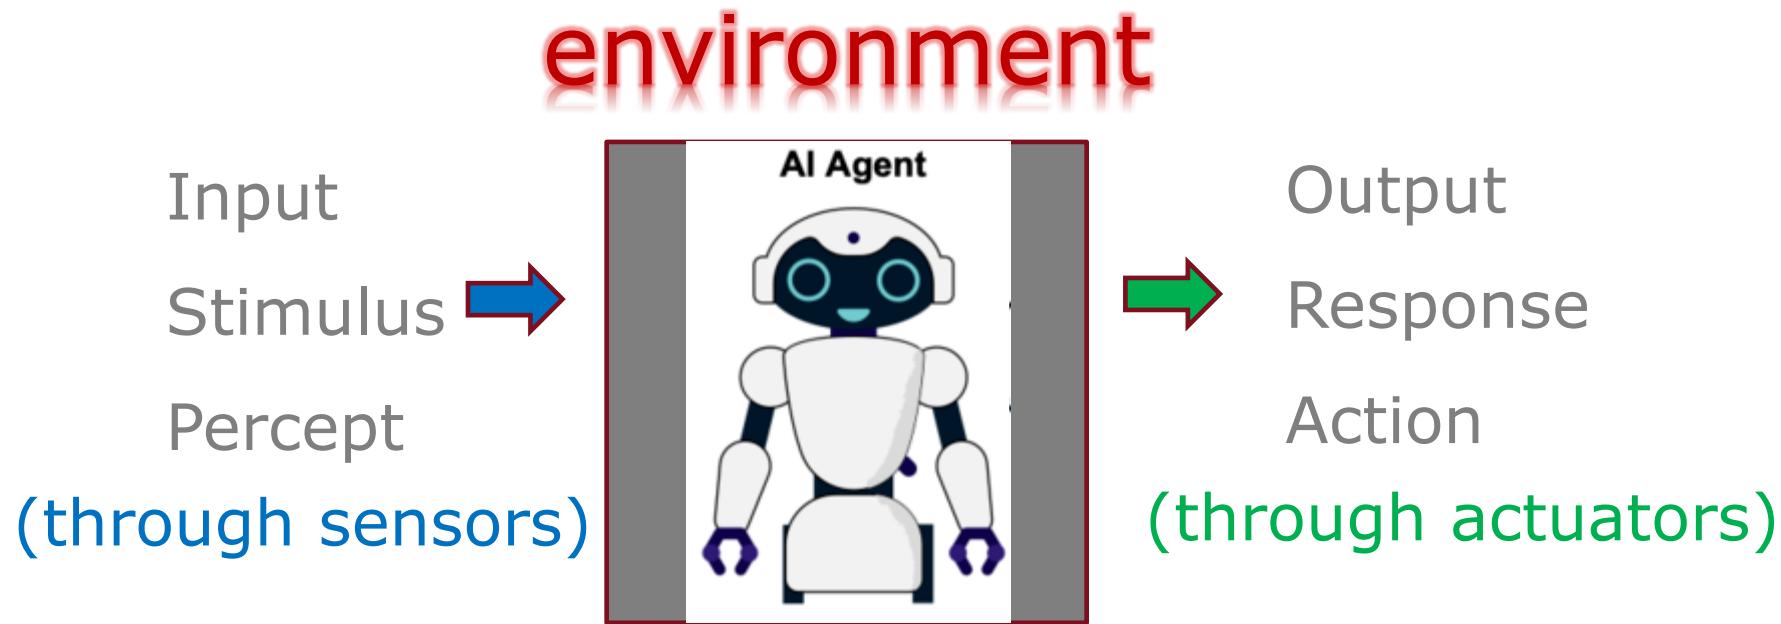


More Specific Example: AI in Games

Game Agent: Interacting with the world around it



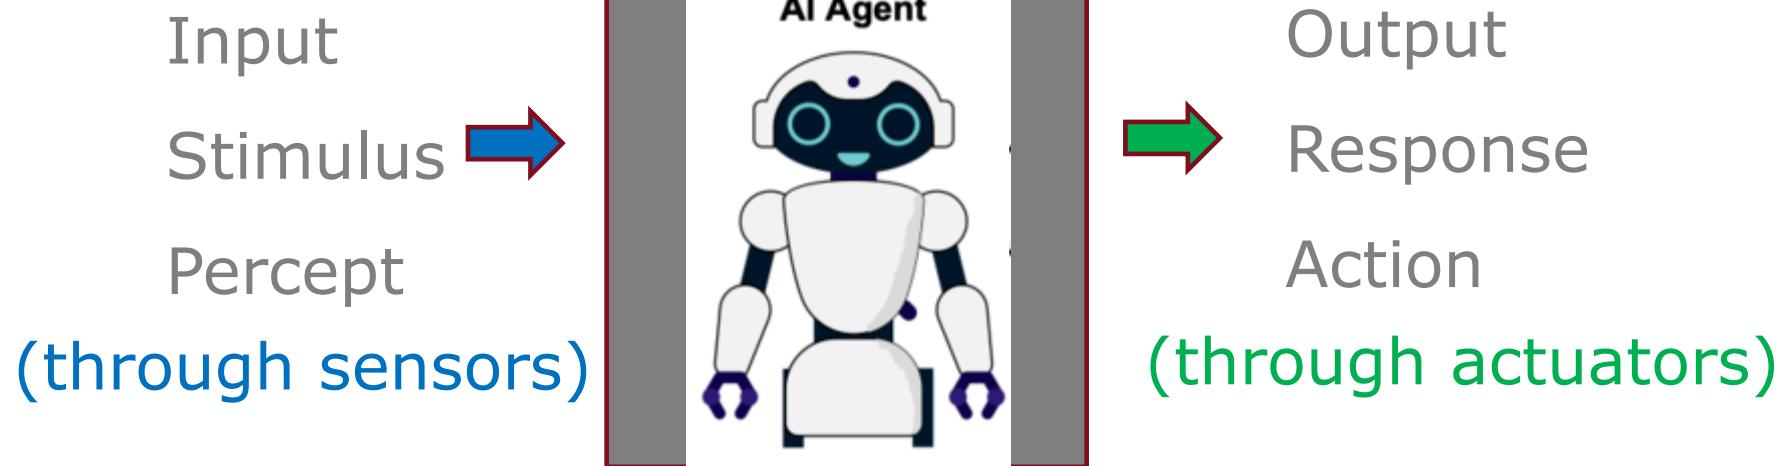
Types of Environments



- An environment could be everything that affects what an intelligent agent perceives and that is affected by the agent's actions.
 - Fully Observable or Partially Observable
 - Deterministic or Non-deterministic (Stochastic)
 - Episodic or Sequential
 - Static or Dynamic
 - Discrete or Continuous
 - Single Agent or Multiple Agents

Fully Observable / Partially Observable

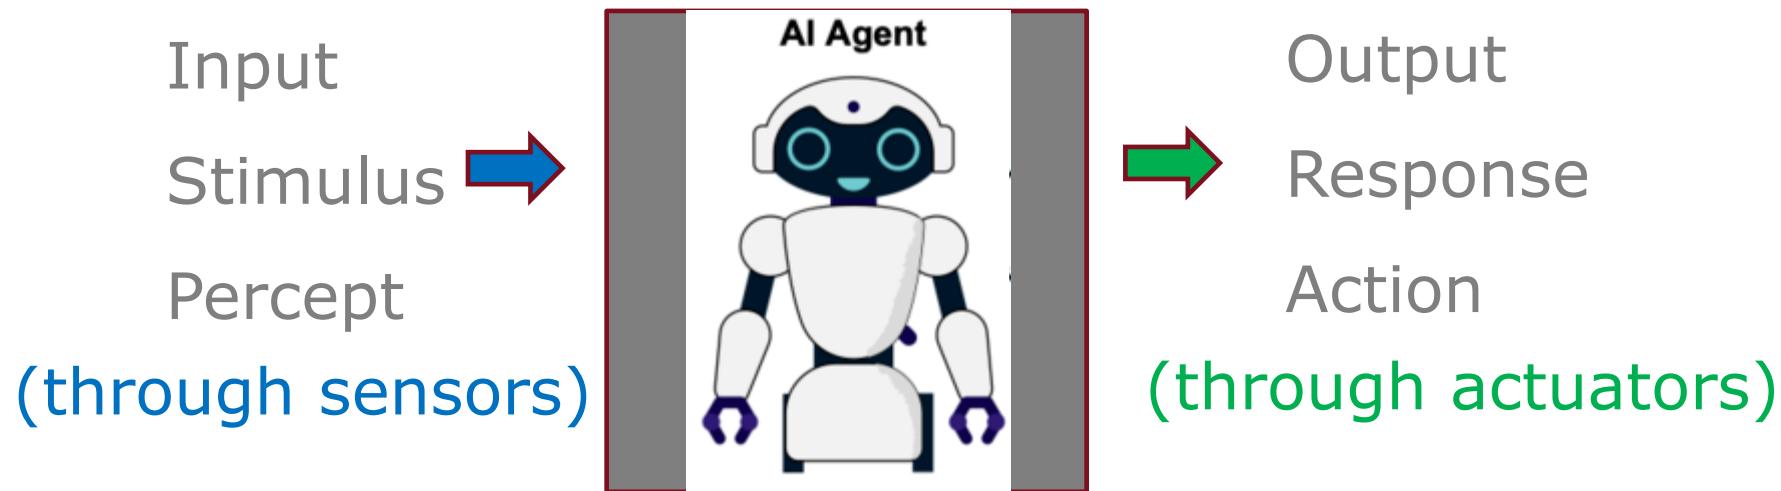
environment



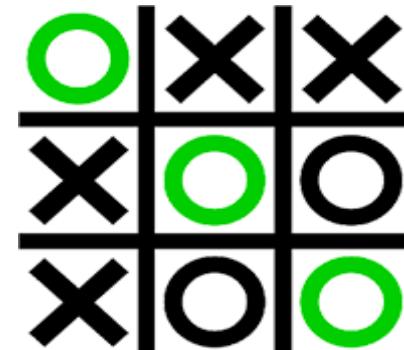
- Fully Observable is that an intelligent agent's sensors give it access to complete state of the environment at each point in time, or else it is partially observable.
- Chess Game's environment is fully observed: An AI player gets to see the entire chess board.
- Poker's environment is partially observed: An AI player gets to see only his own cards and the cards on the table, but not the cards of everyone in the game.



Deterministic / Non-Deterministic (Stochastic) environment

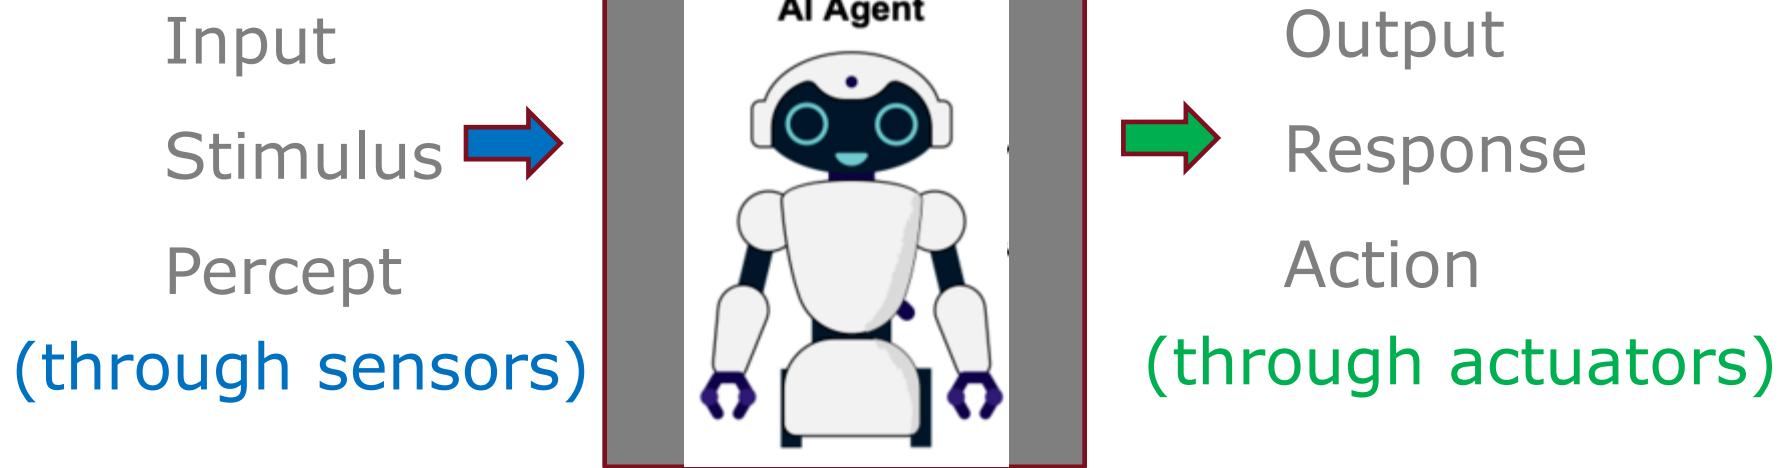


- Deterministic is that if **the next state of the environment** is completely determined by the current state and the next actions of the agent, or else it is non-deterministic (stochastic).
- **Tic Tac Toe's environment** is **Deterministic**
- **Self-driving Car's environment** is **Non-deterministic**



Episodic / Sequential

environment

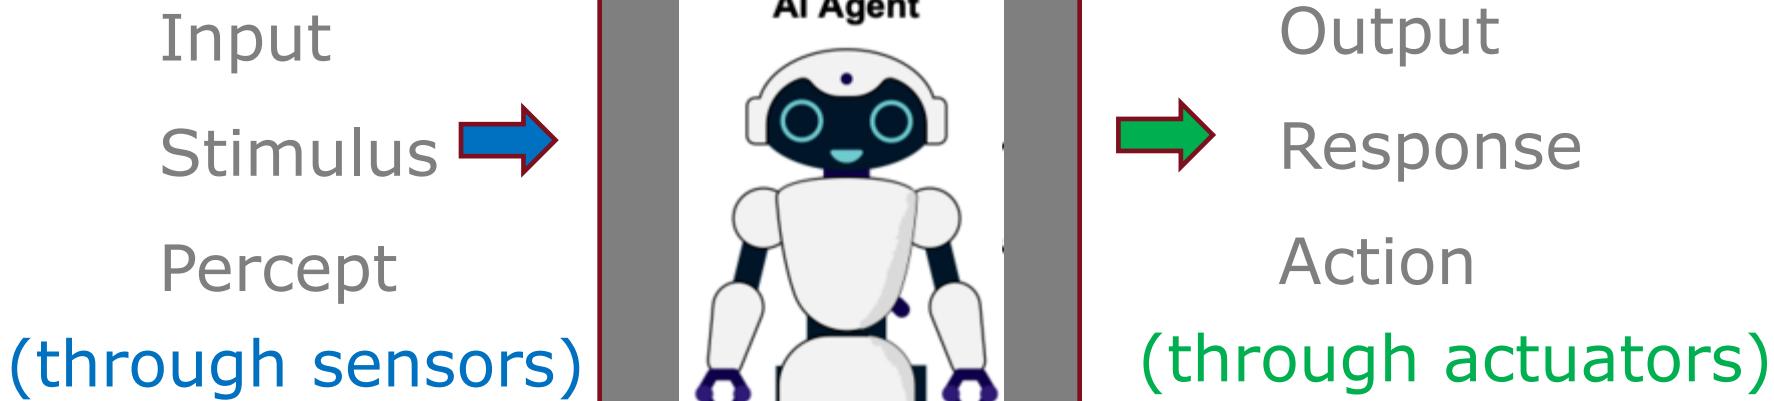


- Episodic: Each single action performed by an intelligent agent depends on only the current atomic episode itself, independent of the previous episode, e.g., **Mail Sorting Machine**
- Sequential: Each next action is affected by the current action, e.g., **Chess Game**



Static / Dynamic

environment



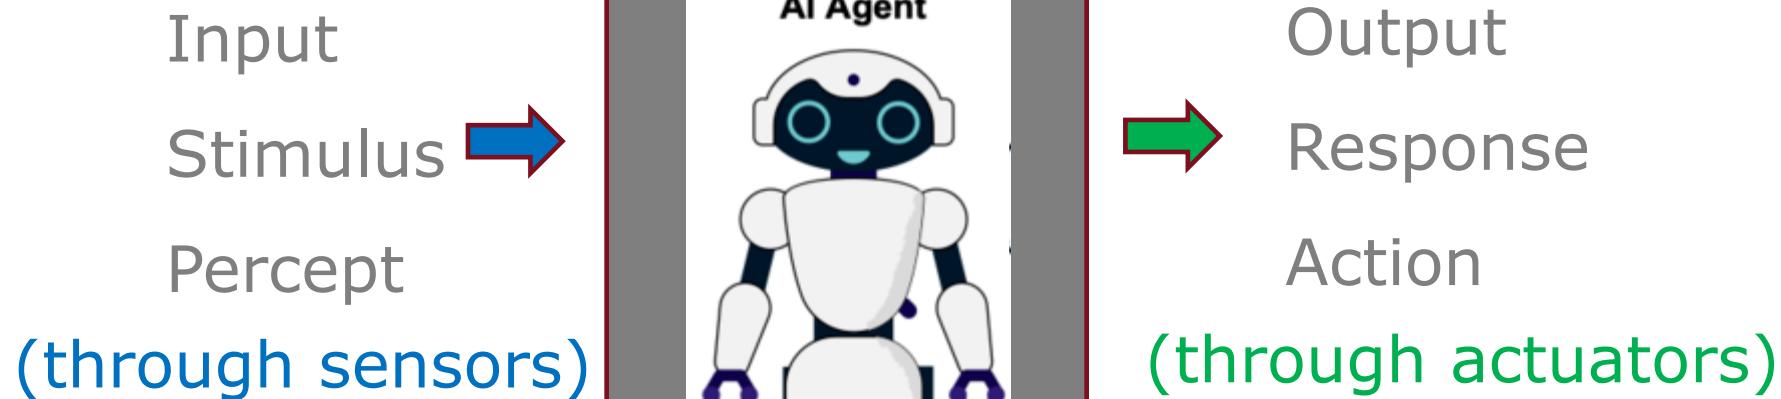
- Static is that if the environment does not change over time constantly while an agent is acting, or else it is dynamic.

- Static: **Crossword Puzzles**
- Dynamic: **Self-driving Taxi.**



Discrete / Continuous

environment

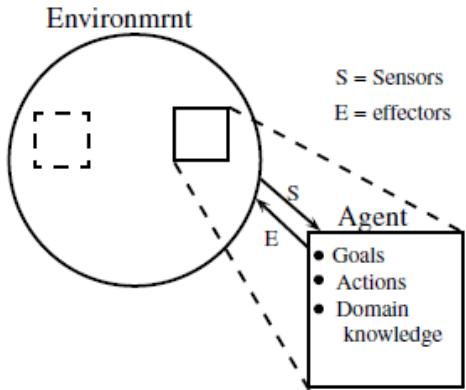


- Discrete: If there are **a limited number of distinct, clearly defined states** of the environment (percepts and actions), the environment is discrete, e.g., **Chess Game** – a set of numbers of discrete/distinct moves of black and white pieces

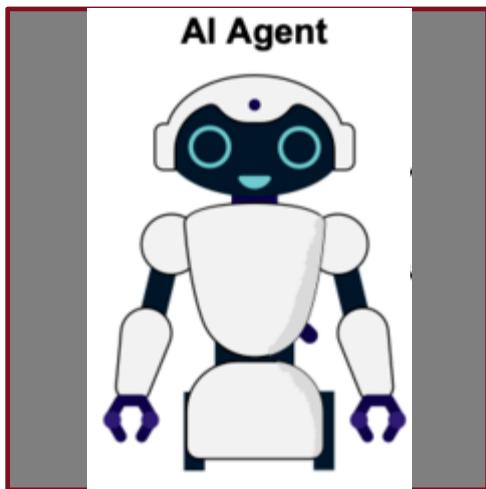


Single Agent / Multiple Agents

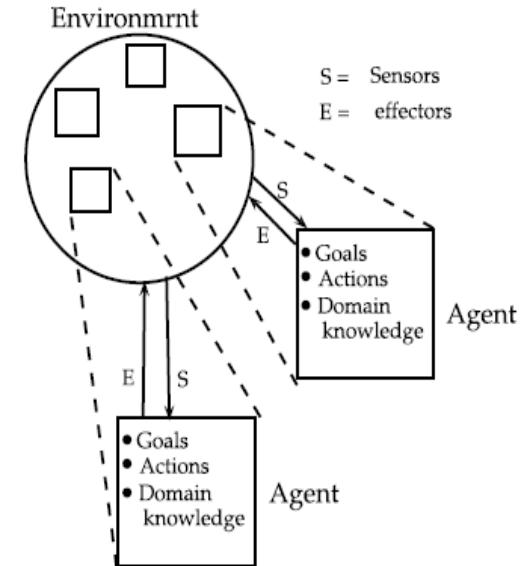
environment



Input
Stimulus →
Percept
(through sensors)



Output
Response
Action
(through actuators)



- An AI agent operating by itself in **an** environment is a **single agent**, e.g., Crossword Puzzle.
- Multi-agents** is when other AI agents are present, e.g., other players in a soccer team or the opposing team.



More AI Agent Examples

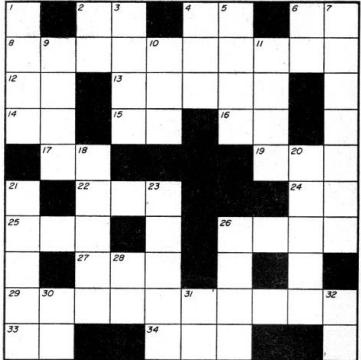
Agent	Sensor	Actuator	Objective	State	Environment
Cleaning Robot	Camera Joint sensor	Limbs Joints	Cleanliness evaluation	Object/joint positions	Cleaned space
Chess agent	Board input interface	Move output interface	Position evaluation	Chess position	Chess board
Self-driving car	Camera Sound sensor	Car control interface	Driving safety and goals	Speed/Position in traffic	Traffic conditions
Chatbot	Keyboard	Screen	Chat evaluation	Dialog history	Chat participants

	Fully observable?	Deterministic?	Episodic?	Static?	Discrete?	Single agent?	Environment	Accessible	Deterministic	Episodic	Static	Discrete
Solitaire	No	Yes	Yes	Yes	Yes	Yes	Chess with a clock Chess without a clock Poker Backgammon Taxi driving Medical diagnosis system Image-analysis system Part-picking robot Refinery controller Interactive English tutor	Yes Yes No Yes No No Yes No No No	Yes Yes No No No No Yes Yes No No	No No No No No No Yes Yes Yes Yes	Semi Yes Yes Yes Yes No Semi Yes Yes No	Yes Yes Yes Yes Yes No No No No Yes
Backgammon	Yes	No	No	Yes	Yes	No						
Taxi driving	No	No	No	No	No	No						
Internet shopping	No	No	No	No	Yes	No						
Medical diagnosis	No	No	No	No	No	Yes						

Task Environments: What Type?

Task Environment	Observable	Agent	Determ.	Episodic	Static	Discrete
Crossword Puzzle	Fully	S	Determ.	Sequen.	Static	Discrete
Chess w. a Clock	Fully	M	Determ.	Sequen.	Semi	Discrete
Poker	Partially	M	Stoch.	Sequen.	Static	Discrete
Backgammon	Fully	M	Stoch.	Sequen.	Static	Discrete
Taxi Driving	Partially	M	Stoch.	Sequen.	Dynamic	Continuous
Medical Diagnostics	Partially	S	Stoch.	Sequen.	Dynamic	Continuous
Image Analysis	Fully	S	Determ.	Episod.	Semi	Continuous
English Tutor	Partially	M	Stoch.	Sequen.	Dynamic	Discrete

Task Environments: More Examples



crossword



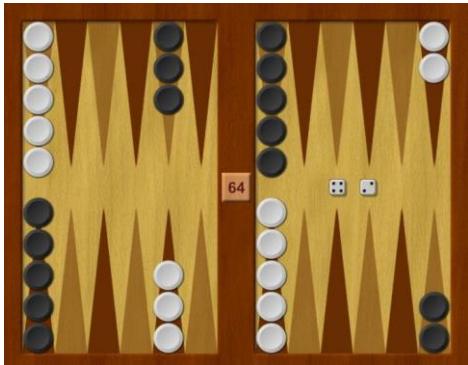
chess with timer



poker



autonomous
taxi



backgammon



medical
diagnostics

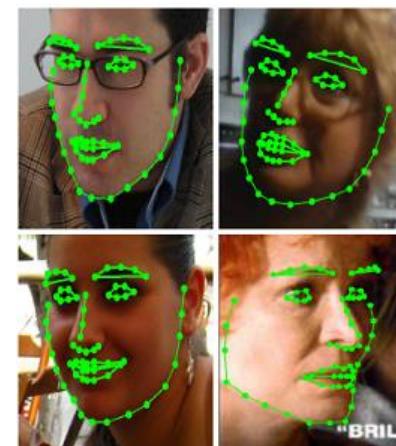


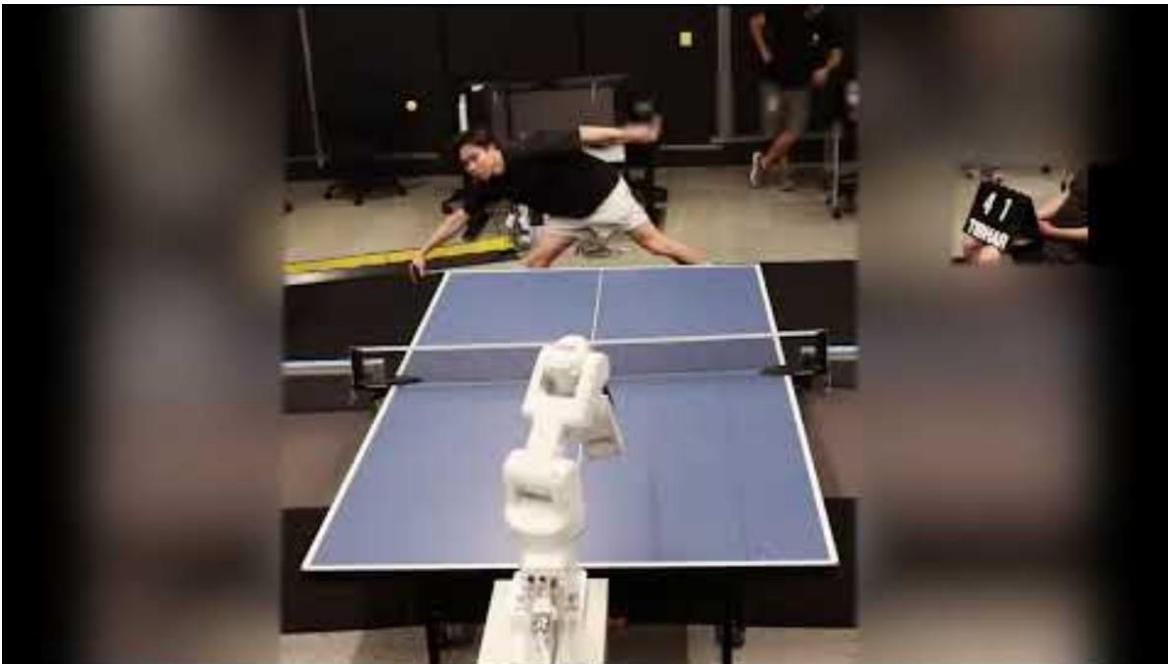
image
processing



interactive
English tutor

Class Exercise 1

- Playing a table tennis match.



- Knitting a sweater.

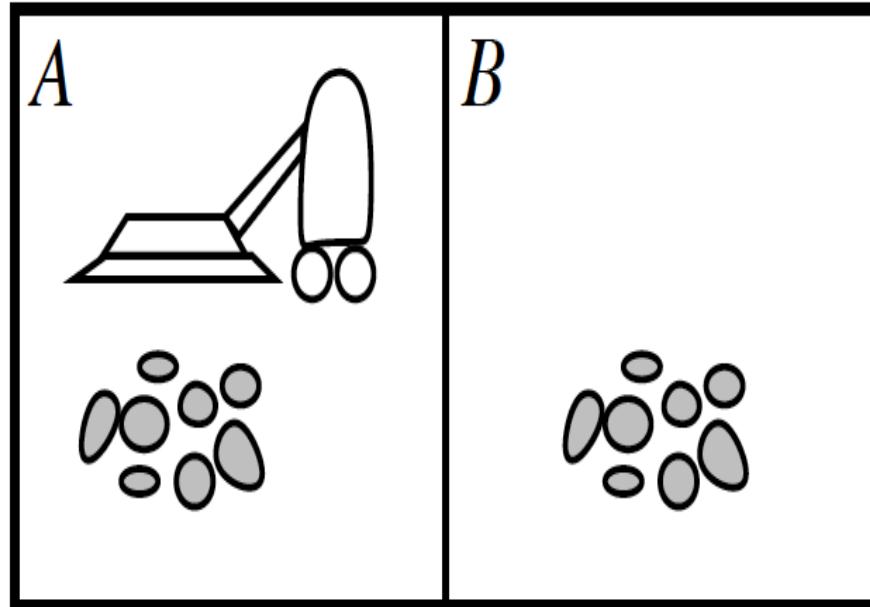


- Fully Observable or Partially Observable
- Deterministic or Non-deterministic (Stochastic)
- Episodic or Sequential
- Static or Dynamic
- Discrete or Continuous
- Single Agent or Multiple Agents

The Vacuum Agent in the Reference Textbook



Vacuum Agent



Percepts: location and contents, e.g., [A, Dirty]

Actions: Left, Right, Suck, NoOp

An agent is **a function F** from percept environments to actions, i.e., $F: P^* \rightarrow A$

Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
:	:

function REFLEX-VACUUM-AGENT([location, status]) returns an action

```
if status = Dirty then return Suck  
else if location = A then return Right  
else if location = B then return Left
```

Left

Define it

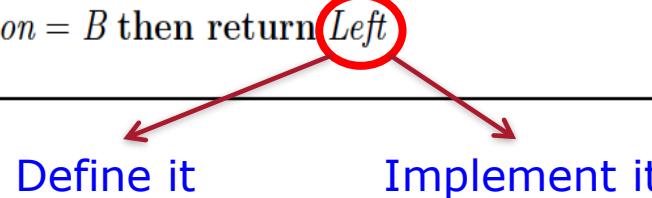
Implement it

Vacuum Agent

Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
:	:

function REFLEX-VACUUM-AGENT([location, status]) returns an action

if *status* = Dirty then return Suck
else if *location* = A then return Right
else if *location* = B then return Left



An agent is **a function *F*** from percept environments to actions, i.e., $F: P^* \rightarrow A$

<https://github.com/aimacode/aima-python/blob/master/agents4e.py>

```
def ReflexVacuumAgent():
    """
```

[Figure 2.8]

A reflex agent for the two-state vacuum environment.

```
>>> agent = ReflexVacuumAgent()
>>> environment = TrivialVacuumEnvironment()
>>> environment.add_thing(agent)
>>> environment.run()
>>> environment.status == {(1,0):'Clean' , (0,0) : 'Clean'}
True
"""
```

```
def program(percept):
    location, status = percept
    if status == 'Dirty':
        return 'Suck'
    elif location == loc_A:
        return 'Right'
    elif location == loc_B:
        return 'Left'

    return Agent(program)
```

Definition of a Rational Agent

Definition: For each possible percept sequence, a **rational agent** should **select the next optimal action that is expected to maximize its performance measure**, i.e., **the expected utility value**, given the evidence provided by the percept sequence and **whatever built-in knowledge the agent has at any given time**.

- A utility function maps a state (or a sequence of states) onto an evaluation value (usually a real number), e.g., cost, profit, score, etc.

Utility × Percepts × Knowledge → Action

- A selected action is optimal if it takes the agent to a state of maximum expected utility given available percepts and knowledge. The agent is rational if it always chooses optimal actions
- Rational is related to **exploration**, **learning**, and **autonomy**

Exploration, Learning, and Autonomy

Exploration: Information gathering in an initially unknown environment



Learning: Updating the knowledge about the environment during perception

Exploration + Learning:

- Maximizing the expected performance
- Doing actions in order to modify future percepts



Autonomy: Extend to which an agent relies on the prior knowledge of its designer – lacks autonomy.



Task Environments

What is the **first step** in designing a **rational AI agent**? That is to specify the task environment.

The task environment is defined using **PEAS** descriptors:

- **Performance measure**
- **Environment**
- **Actuators**
- **Sensors**

Example: Pepper



Performance Measure	Number of correctly answered questions
Environment	Rigi Kaltbad Station Shop
Actuators	loud speaker
Sensors	camera, microphones

Example: Internet Shopping Agent

Performance measures: price, quality, appropriateness, efficiency

Environment: current and future WWW sites, vendors, shippers

Actuators: display to user, follow URL, fill in form

Sensors: HTML pages (text, graphics, scripts)



Example: Internet Shopping Agent



Example: AI Healthcare Agent

- **Performance measure:** Healthy patient, minimize costs, lawsuits
- **Environment:** Patient, hospital, staff
- **Actuators:** Screen display (questions, tests, diagnoses, treatments, referrals)
-
- **Sensors:** Keyboard (entry of symptoms, findings, patient's answers)



Example: AI Healthcare Agent



Example: Automated Taxi

Performance measures: safety, destination, profits, legality, comfort, etc.

Environment: US streets/freeways, traffic, pedestrians, customers, weather, etc.

Actuators: steering, accelerator, brake, horn, speakers, display, etc.

Sensors: cameras, accelerometers, gauges, engine sensors, keyboard, GPS, etc.



Waymo One: Google-launched autonomous taxi service

Example: Automated Taxi



Other PEAS Examples

Agent Type	Performance Measure	Environment	Actuators	Sensors
Medical diagnosis system	healthy patient, costs, lawsuits	patient, hospital, stuff	display questions, tests, diagnoses, treatments, referrals	keyboard entry of symptoms, findings, patient's answers
Satellite image analysis system	correct image categorization	downlink from orbiting satellite	display categorization of scene	color pixel arrays
Part-picking robot	percentage of parts in correct bins	conveyor belt with parts, bins	jointed arm and hand	camera, joint angle sensors
Refinery controller	purity, yield, safety	refinery, operators	valves pumps, heaters displays	temperature, pressure, chemical sensors
Interactive English tutor	student's score on test	set of students, testing agency	display exercises, suggestions, corrections	keyboard entry

Class Exercise 2

- AI-Powered Microwave Oven.



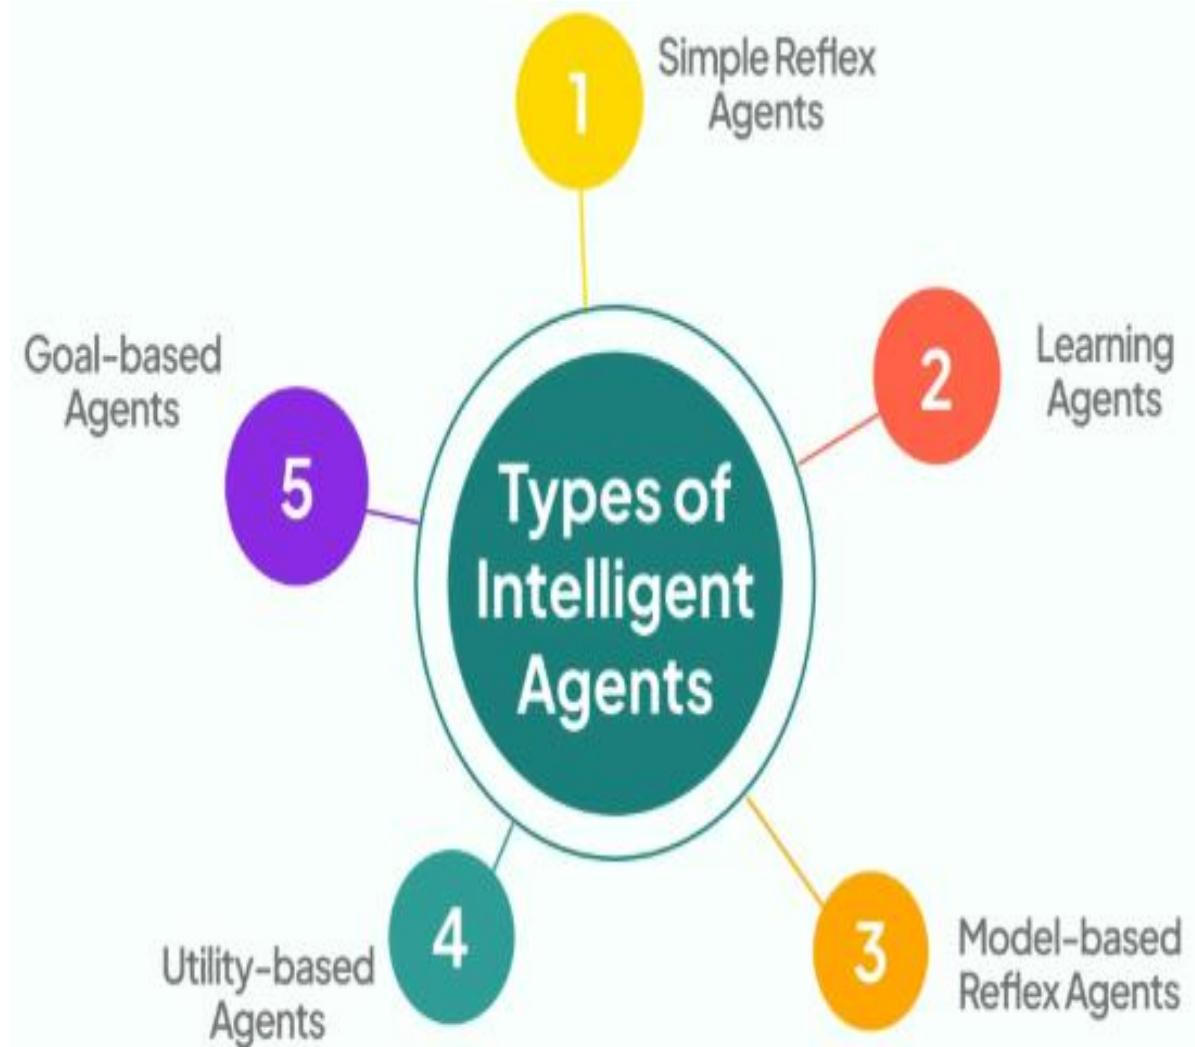
- Autonomous Supply Delivery Plane.



- Performance Measure
- Sensors
- Actuators
- Environment

Agents: Types of Agents (Programs)

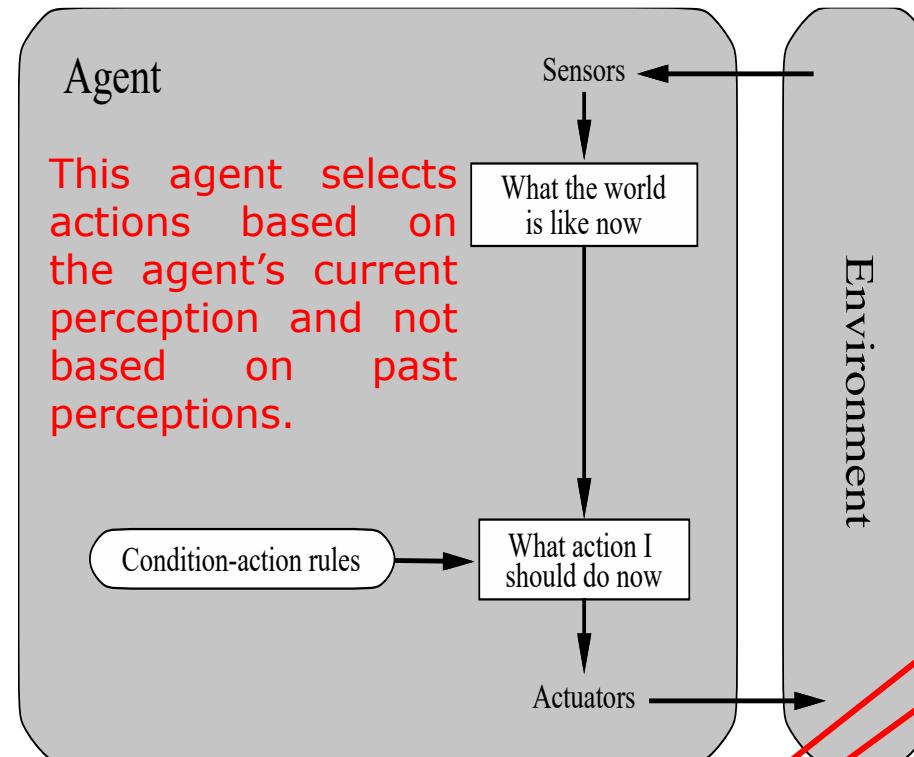
- **Simple Reflex (Reactive) Agents:** an agent whose action depends only on the current percept.
- **Model-based Reflex Agents:** an agent whose action is derived directly from an internal model of the current world state that is updated over time.
- **Goal-based Agents:** an agent that selects actions that it believes will achieve explicitly represented goals.
- **Utility-based Agents:** an agent that selects actions that it believes will maximize the expected utility of the outcome state.
- **Learning Agents:** an agent whose behavior/action improves over time based on its experience.



<https://github.com/aimacode/aima-python/blob/master/agents4e.py>

Simple Reflex (Reactive) Agent

- Senses the world and responds immediately based on the **simple rules whose condition matches the current state** based on percept and then acts accordingly.
- No explicit world model, no "memory"**
- This agent will only work if the environment is fully observable.
- Agent: Mail sorting robot
- Environment: Conveyor belt of letters
- Rule: e.g. if city = Worcester, then put Massachusetts bag



```
function SIMPLE-REFLEX-AGENT(percept) returns an action
  persistent: rules, a set of condition-action rules
  state  $\leftarrow$  INTERPRET-INPUT(percept)
  rule  $\leftarrow$  RULE-MATCH(state, rules)
  action  $\leftarrow$  rule.ACTION
  return action
```

INTERPRET-INPUT function generates an abstracted description of the **current state** from the percept.

RULE-MATCH function returns the **1st rule in the set of rules** that matches the given state description.

```
def SimpleReflexAgentProgram(rules, interpret_input):
    """
    [Figure 2.10]
    This agent takes action based solely on the percept.
    """


```

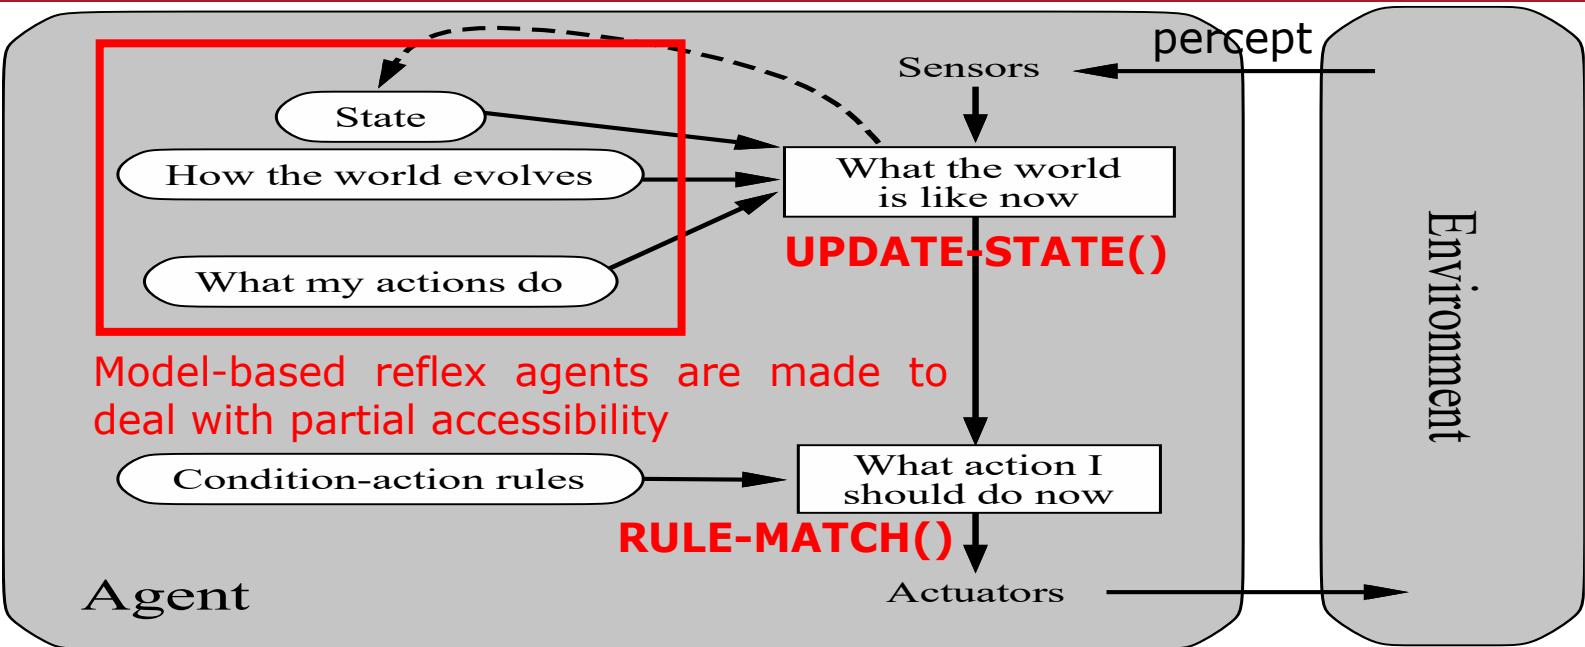
```
def program(percept):
    state = interpret_input(percept)
    rule = rule_match(state, rules)
    action = rule.action
    return action
```

```
return program
```

```
def rule_match(state, rules):
    """
    Find the first rule that matches state.
    """
    for rule in rules:
        if rule.matches(state):
            return rule
```

Model-based Reflex Agents

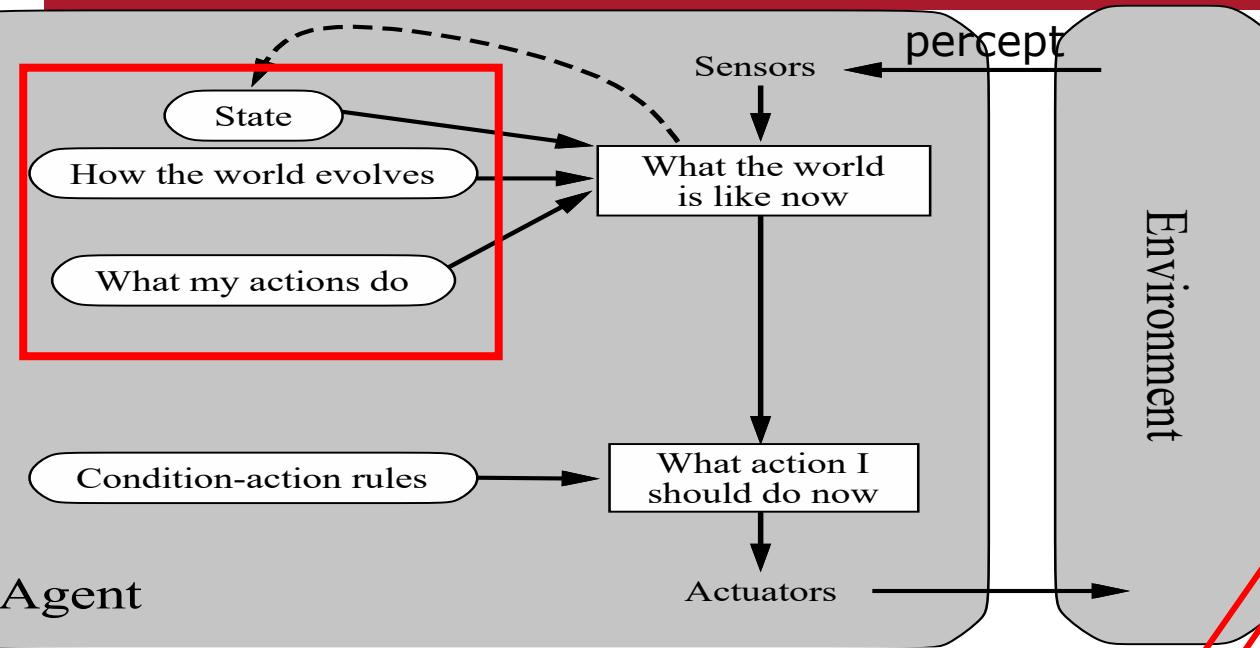
- Keep track of **the previous state** of the world that the agent cannot perceive right now, i.e., **the internal world model (agent state) that depends on what it has seen before, so it holds information on the unobserved aspects of the current state.**
- Uncertainty about the world state is unavoidable because of limited sensing capabilities and the limited world model that represents **the agent's "best guess" of the world state, the evolution of the world and the effects of its actions**
- New Current State = Previous State + Current Percept
- Action may depend on the new current state.



```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
            transition-model, a description of how the next state depends on
            the current state and action
            sensor-model, a description of how the current world state is reflected
            in the agent's percepts
            rules, a set of condition-action rules
            action, the most recent action, initially none

  state ← UPDATE-STATE(state, action, percept, transition-model, sensor-model)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```

Model-based Reflex Agents



```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
  transition-model, a description of how the next state depends on
    the current state and action
  sensor-model, a description of how the current world state is reflected
    in the agent's percepts
  rules, a set of condition-action rules
  action, the most recent action, initially none
```

```
state ← UPDATE-STATE(state, action, percept, transition-model, sensor-model)
rule ← RULE-MATCH(state, rules)
action ← rule.ACTION
return action
```

```
def ModelBasedReflexAgentProgram(rules, update_state, transition_model, sensor_model):
    """
    [Figure 2.12]
    This agent takes action based on the percept and state.
    """

    def program(percept):
        program.state = update_state(program.state, program.action, percept, transition_model, sensor_model)
        rule = rule_match(program.state, rules)
        action = rule.action
        return action

    program.state = program.action = None
    return program

def rule_match(state, rules):
    """
    Find the first rule that matches state.
    """
    for rule in rules:
        if rule.matches(state):
            return rule
```

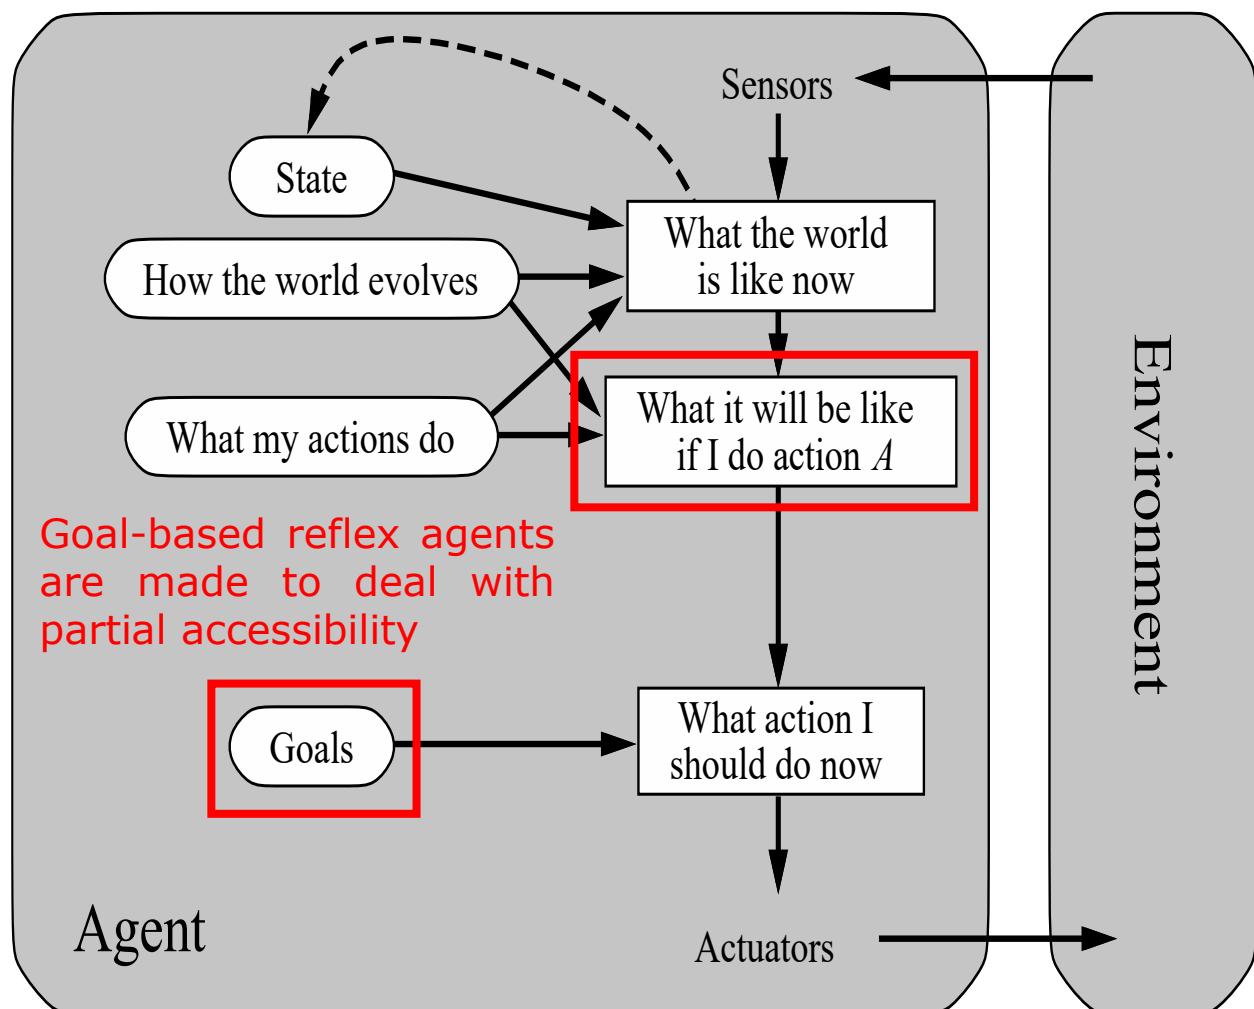
RULE-MATCH function
returns the 1st rule in the set of rules that matches the given state description.

Goal-based Reflex Agents

- Builds a model of the world and uses an explicit representation of goals
- Evaluates the effects/impacts of actions on the world model before selecting an action**
- This agent **reviews many actions** and chooses the **first** available one which can achieve its goals

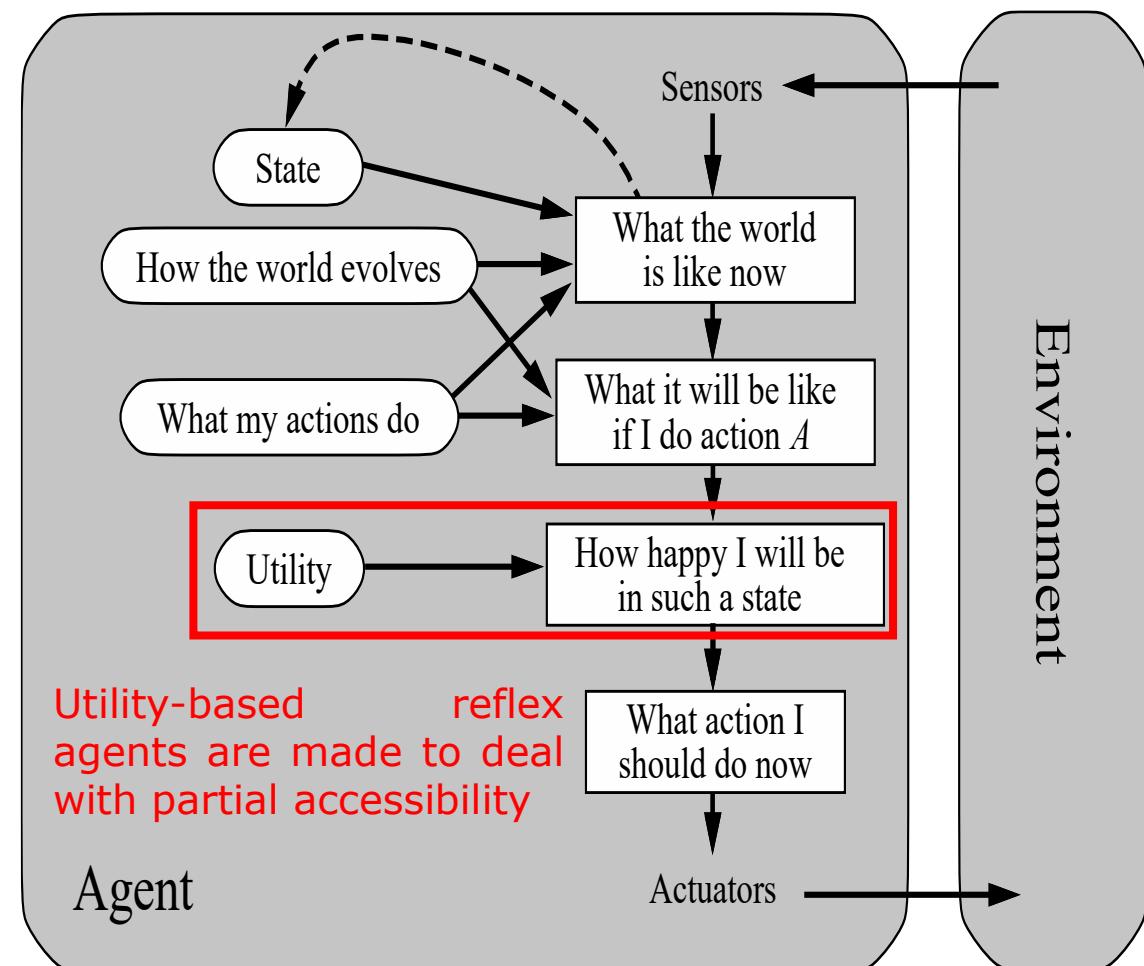
Goal based agent = Goal information + Model-based Agent

To choose action that achieve the goal



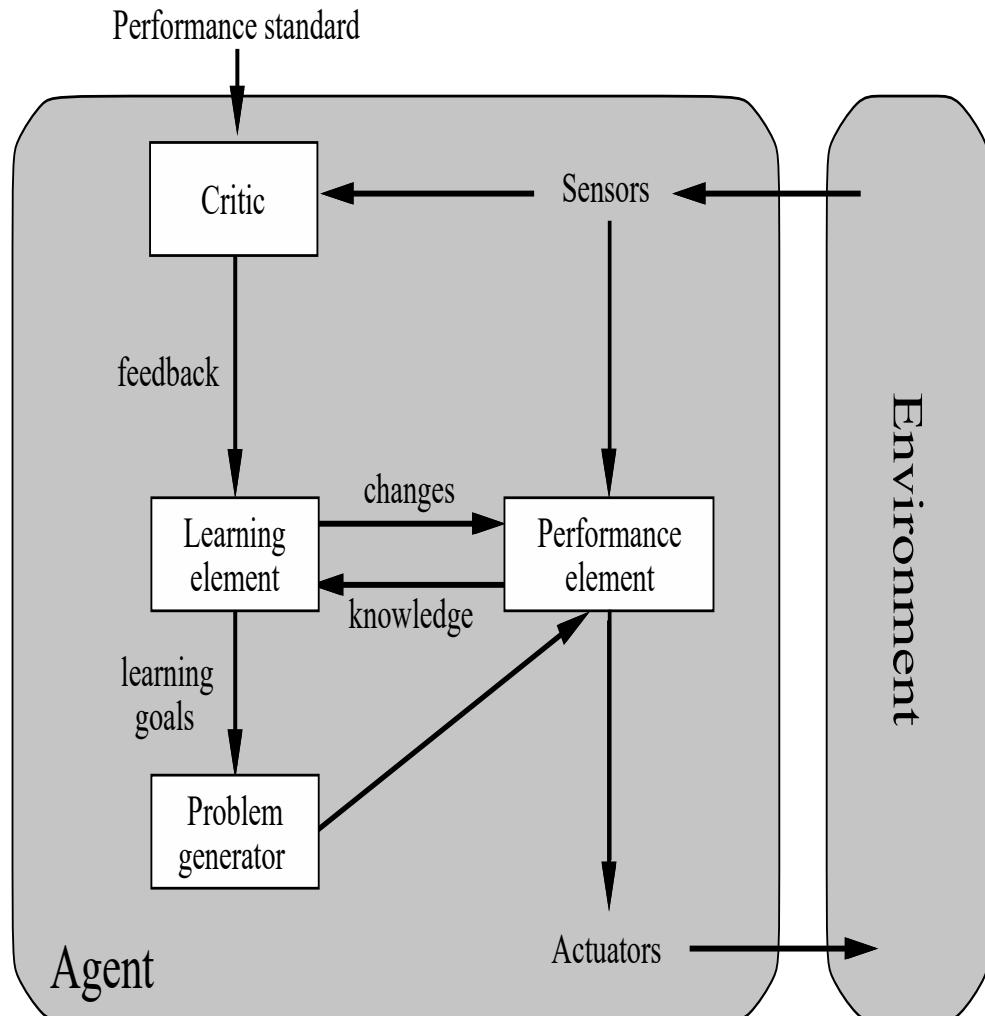
Utility-based Agents

- We may have several actions which all satisfy our goal, so we need some way of working out the most efficient one.
- A utility function maps each state after each action to a real number representing how **efficiently** each action achieves the goal.
- A goal-based navigation agent is tasked with getting from point A to point B. If the agent succeeds, the goal has been satisfied.
- A utility-based navigation agent could seek to get from point A to point B in the shortest amount of time, with the minimum expenditure of fuel, or both.

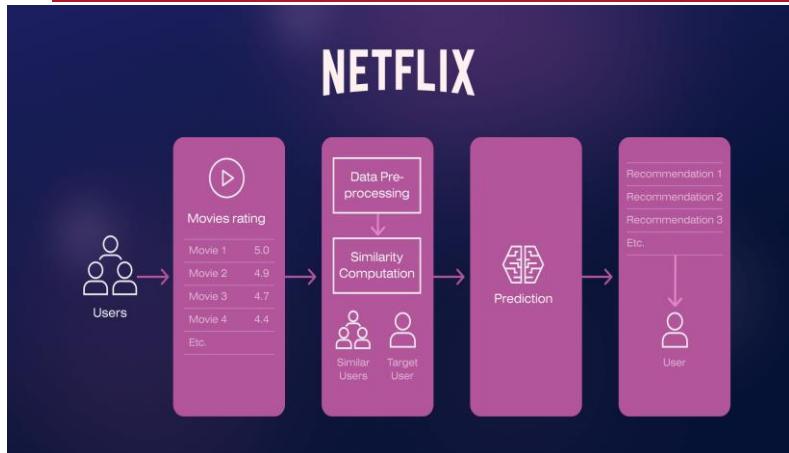


Learning Agents

- This agent can acquire new skills and reflect on its own performance to improve over time.
- Consists of four basic components: Critic, Learning Element, Performance Element, Problem Generator
- **Performance element:** select external actions to act on the environment
- **Critic:** evaluates the behavior of the agent based on its performance and gives the evaluation as feedback to the learning element
- **Learning element:** improves the performance element by posing new tasks
- **Problem generator:** suggests actions that will lead to new and informative experiences



Class Exercise 3



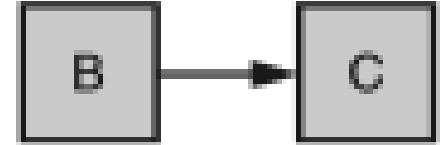
Simple Reflex (Reactive) Agents
Model-based Reflex Agents
Goal-based Agents
Utility-based Agents
Learning Agents



How Can We Represent The State of Environment That The Agent Inhibits?

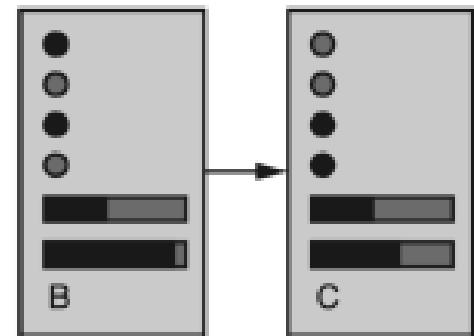
- Atomic Representation

- each state of the world is indivisible (a black box with no internal structure)
 - used in search and game-playing algorithms, Markov Decision Processes



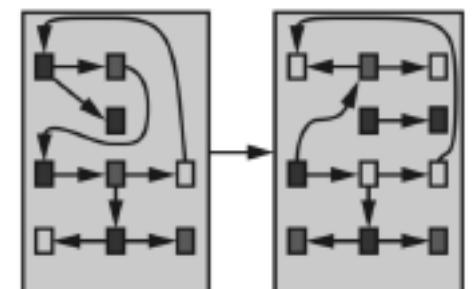
- Factored Representation

- each state splits into a fixed set of **variables (attributes)**, each of which can have a value, e.g., Boolean, real-valued, or one of a fixed set of symbols.
 - used in constraint satisfaction, propositional logic, and planning



- Structured Representation

- each state consists of a set of objects (each may have attributes) with various and varying **relationships**
 - used in first-order logic



Summary

- Agents interact with environments through **actuators** and **sensors**
- The **agent function** describes what the agent does in all circumstances
- The **expected performance measure** evaluates the environment sequence
- A perfectly **rational** agent maximizes expected performance
- **Agent programs** implement (some) agent functions
- **PEAS** descriptions define task environments
- **Environments** are categorized along several dimensions: observable; deterministic; episodic; static; discrete; single-agent
- Several basic **agent architectures** exist: simple reflex, model-based reflex, goal-based, utility-based, and learning-based.